

# Recent Advances in Underwater Basket Weaving Under the Extreme Pressure of the Mariana Trench

André Lauren Benjamin<sup>1</sup>, Calvin Cordozar Broadus Jr.<sup>2,3</sup> (✉), and Antwan André Patton<sup>1</sup>[0000–1111–2222–3333]

<sup>1</sup> Fictional Southern University, Savannah GA 31404, USA  
[{a.l.benjamin,a.a.patton}@fsu.fake](mailto:{a.l.benjamin,a.a.patton}@fsu.fake)

<sup>2</sup> Fictional West Coast University, Long Beach CA 90840, USA [ccb@fwcu.fake](mailto:ccb@fwcu.fake)

<sup>3</sup> Secondary European Affiliation, Tiergartenstr. 17, 69121 Heidelberg, Germany  
[lncs@springer.com](mailto:lncs@springer.com)

**Abstract.** This document provides a basic paper template and submission guidelines. Abstracts must be a single paragraph, ideally between 4–6 sentences long. Gross violations will trigger corrections at the camera-ready phase.

**Keywords:** First keyword · Second keyword · Another keyword.

## 1 Related Literature

### 1.1 Background on Counterfactual Explanations

[11, 5, 2]

### 1.2 Learning Representations

For example, joint-energy models

### 1.3 Generalization and Robustness

[9] generate counterfactual images for MNIST and ImageNet through independent mechanisms (IM): each IM learns class-conditional input distributions over a specific lower-dimensional, semantically meaningful factor, such as *texture*, *shape* and *background*. They demonstrate that using these generated counterfactuals during classifier training improves model robustness. Similarly, [1] argue that counterfactuals represent potentially useful training data in machine learning, especially in supervised settings where inputs may be reasonably mapped to multiple outputs. They, too, demonstrate that augmenting the training data of image classifiers can improve generalization.

[10] propose an approach using counterfactuals in training that does not rely on data augmentation: they argue that counterfactual pairs typically already

exist in training datasets. Specifically, their approach relies on, firstly, identifying similar input samples with different annotations and, secondly, ensuring that the gradient of the classifier aligns with the vector between pairs of counterfactual inputs using the cosine distance as a loss function (referred to as *gradient supervision*) (*this might be useful for our task as well*). In the natural language processing (NLP) domain, counterfactuals have similarly been used to improve models through data augmentation: [12], propose POLYJUICE, a general-purpose counterfactual generator for language models. They demonstrate empirically that augmenting training data through POLYJUICE counterfactuals improves robustness in a number of NLP tasks.

#### 1.4 Link to Adversarial Training

[3] propose two definitional differences between Adversarial Examples (AE) and Counterfactual Explanations (CE): firstly, and more importantly according to the authors, the term AE implies missclassification, which is not the case for CE (*this might be a useful notion for use to distinguish between adversarials and explanations during training*); secondly, they argue that closeness plays a more critical role in the context of CE but confess that even counterfactuals that are not close might be relevant explanations. [7] show that CE and AE are equivalent under certain conditions and derive upper bounds on the distances between them.

#### 1.5 Closely Related

[4] are the first to propose end-to-end training pipeline that includes counterfactual explanations as part of the training procedure. In particular, they propose a specific network architecture that includes a predictor and CE generator network (*akin a GAN?*), where the parameters of the CE generator network are learnable. Counterfactuals are generated during each training iteration and fed back to the predictor network (*here we are aligned*). In contrast, we impose no restrictions on the neural network architecture at all. (*to ensure the one-hot encoding of categorical features is maintained, they simple use softmax (might be interesting for CE.jl)*) Interestingly, the authors find that their approach is sensitive to the choice of the loss function: only MSE seems to lead to good performance. They also demonstrate theoretically, that the objective function is difficult to optimize due to divergent gradients and suffers from poor adversarial robustness. (*because partial gradients with respect to the classification loss component and the counterfactual validity component point in opposite directions*). To mitigate these issues, the authors use block-wise gradient descent: they first update with respect to classification loss and then use a second update with respect to the other loss components (*this might be useful for our task as well*). [8] propose a way to train models that are guaranteed to provide recourse for individuals with high probability. The approach builds on adversarial training (*here we are aligned*), where in this context adversarial examples are actively encouraged to exist, but

only target attacks with respect to the positive class. The proposed method allows for imposing a set of actionable recourse ex-ante: for example, users can impose mutability constraints for features (*here we are aligned*). (*To solve their objective function more efficiently, they use a first-order Taylor approximation to approximate the recourse loss component (might be applicable in our case)*)

[6] introduce Counterfactual Adversarial Training (CAT) with intention of improving generalization and robustness of language models. Specifically, they propose to proceed as follows: firstly, identify training samples that are subject to high predictive uncertainty (entropy); secondly, generate counterfactual explanations for those samples; and, finally, finetune the model on the augmented dataset that includes the generated counterfactuals.

**Acknowledgments.** A bold run-in heading in small font size at the end of the paper is used for general acknowledgments, for example: This study was funded by X (grant number Y).

**Disclosure of Interests.** It is now necessary to declare any competing interests or to specifically state that the authors have no competing interests. Please place the statement with a bold run-in heading in small font size beneath the (optional) acknowledgments, for example: The authors have no competing interests to declare that are relevant to the content of this article. Or: Author A has received research grants from Company W. Author B has received a speaker honorarium from Company X and owns stock in Company Y. Author C is a member of committee Z.

## Bibliography

- [1] Abbasnejad, E., Teney, D., Parvaneh, A., Shi, J., van den Hengel, A.: Counterfactual vision and language learning. In: 2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR). pp. 10041–10051 (2020). <https://doi.org/10.1109/CVPR42600.2020.01006>
- [2] Altmeyer, P., Farmanbar, M., van Deursen, A., Liem, C.C.: Faithful model explanations through energy-constrained conformal counterfactuals. In: Proceedings of the AAAI Conference on Artificial Intelligence. vol. 38, pp. 10829–10837 (2024)
- [3] Friesleben, T.: The intriguing relation between counterfactual explanations and adversarial examples. *Minds and Machines* **32**(1), 77–109 (2022)
- [4] Guo, H., Nguyen, T.H., Yadav, A.: Counternet: End-to-end training of prediction aware counterfactual explanations. In: Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining. p. 577–589. KDD ’23, Association for Computing Machinery, New York, NY, USA (2023). <https://doi.org/10.1145/3580305.3599290>, <https://doi.org/10.1145/3580305.3599290>
- [5] Joshi, S., Koyejo, O., Vigitbenjaronk, W., Kim, B., Ghosh, J.: Towards realistic individual recourse and actionable explanations in black-box decision making systems (2019)
- [6] Luu, H.L., Inoue, N.: Counterfactual adversarial training for improving robustness of pre-trained language models. In: Proceedings of the 37th Pacific Asia Conference on Language, Information and Computation. pp. 881–888 (2023)
- [7] Pawelczyk, M., Agarwal, C., Joshi, S., Upadhyay, S., Lakkaraju, H.: Exploring counterfactual explanations through the lens of adversarial examples: A theoretical and empirical analysis. In: Camps-Valls, G., Ruiz, F.J.R., Valera, I. (eds.) Proceedings of The 25th International Conference on Artificial Intelligence and Statistics. Proceedings of Machine Learning Research, vol. 151, pp. 4574–4594. PMLR (28–30 Mar 2022), <https://proceedings.mlr.press/v151/pawelczyk22a.html>
- [8] Ross, A., Lakkaraju, H., Bastani, O.: Learning models for actionable recourse. In: Proceedings of the 35th International Conference on Neural Information Processing Systems. NIPS ’21, Curran Associates Inc., Red Hook, NY, USA (2024)
- [9] Sauer, A., Geiger, A.: Counterfactual generative networks (2021), <https://arxiv.org/abs/2101.06046>
- [10] Teney, D., Abbasnejad, E., van den Hengel, A.: Learning what makes a difference from counterfactual examples and gradient supervision. In: Computer Vision–ECCV 2020: 16th European Conference, Glasgow, UK, August 23–28, 2020, Proceedings, Part X 16. pp. 580–599. Springer (2020)

- [11] Wachter, S., Mittelstadt, B., Russell, C.: Counterfactual explanations without opening the black box: Automated decisions and the GDPR. *Harv. JL & Tech.* **31**, 841 (2017). <https://doi.org/10.2139/ssrn.3063289>
- [12] Wu, T., Ribeiro, M.T., Heer, J., Weld, D.: Polyjuice: Generating counterfactuals for explaining, evaluating, and improving models. In: Zong, C., Xia, F., Li, W., Navigli, R. (eds.) *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*. pp. 6707–6723. Association for Computational Linguistics, Online (Aug 2021). <https://doi.org/10.18653/v1/2021.acl-long.523>, <https://aclanthology.org/2021.acl-long.523>

## 2 Appendix

### 2.1 Initial Grid Search

#### Generator Params

*Linearly Separable*

*Moons*

*Circles*

**Table 1.** Results for Linearly Separable data by energy penalty.

Objective	Lambda	Energy (exper)	Generator	Type	Value	Std
full		0.5	<i>ECCo</i>		-11.7601	9.83205
full		0.5	<i>Generic</i>		-1.064e18	2.41995e18
<b>full</b>		<b>0.5</b>	<b>Omniscient</b>		<b>-2.5405</b>	<b>0.122789</b>
full		0.5	<i>REVISE</i>		-15.0277	12.5588
vanilla		0.5	<i>ECCo</i>		-4.39923	3.65268
vanilla		0.5	<i>Generic</i>		-4.38184	3.48393
vanilla		0.5	<i>Omniscient</i>		-5.24831	4.61237
vanilla		0.5	<i>REVISE</i>		-4.94731	4.2233
full		1.0	<i>ECCo</i>		-11.5401	11.0622
full		1.0	<i>Generic</i>		-1.70667e11	3.88205e11
<b>full</b>		<b>1.0</b>	<b>Omniscient</b>		<b>-2.58956</b>	<b>0.117255</b>
full		1.0	<i>REVISE</i>		-15.7258	13.2676
vanilla		1.0	<i>ECCo</i>		-4.27742	3.50817
vanilla		1.0	<i>Generic</i>		-4.44409	3.4741
vanilla		1.0	<i>Omniscient</i>		-5.11353	4.4628
vanilla		1.0	<i>REVISE</i>		-4.91409	4.24885
full		5.0	<i>ECCo</i>		-3.99166	3.12284
full		5.0	<i>Generic</i>		-4.88333e17	1.11064e18
<b>full</b>		<b>5.0</b>	<b>Omniscient</b>		<b>-2.5325</b>	<b>0.117196</b>
full		5.0	<i>REVISE</i>		-14.5887	12.1265
vanilla		5.0	<i>ECCo</i>		-4.39614	3.64978
vanilla		5.0	<i>Generic</i>		-4.37909	3.48341
vanilla		5.0	<i>Omniscient</i>		-5.24668	4.60676
vanilla		5.0	<i>REVISE</i>		-4.94655	4.22198
<b>full</b>		<b>10.0</b>	<b>ECCo</b>		<b>-2.30721</b>	<b>0.73475</b>
full		10.0	<i>Generic</i>		-1.69667e11	3.85893e11
full		10.0	<i>Omniscient</i>		-2.53433	0.116736
full		10.0	<i>REVISE</i>		-15.5346	13.0245
vanilla		10.0	<i>ECCo</i>		-4.28116	3.50992
vanilla		10.0	<i>Generic</i>		-4.4428	3.47049
vanilla		10.0	<i>Omniscient</i>		-5.11933	4.46099
vanilla		10.0	<i>REVISE</i>		-4.91285	4.24407
<b>full</b>		<b>15.0</b>	<b>ECCo</b>		<b>-2.00576</b>	<b>0.48751</b>
full		15.0	<i>Generic</i>		-4.91e17	1.11683e18
full		15.0	<i>Omniscient</i>		-2.52833	0.11602
full		15.0	<i>REVISE</i>		-14.3763	11.7494
vanilla		15.0	<i>ECCo</i>		-4.3957	3.65194
vanilla		15.0	<i>Generic</i>		-4.38497	3.48359
vanilla		15.0	<i>Omniscient</i>		-5.24893	4.60484
vanilla		15.0	<i>REVISE</i>		-4.94518	4.22746

**Table 2.** Results for Moons data by energy penalty.

Objective	Lambda	Energy (exper)	Generator	Type	Value	Std
full		0.5	<i>ECCo</i>	-7.08577	7.51393	
full		0.5	<i>Generic</i>	-1.1064e31	2.53239e31	
<b>full</b>	<b>0.5</b>	<b>Omniscient</b>	<b>-4.58057</b>	<b>4.8256</b>		
full		0.5	<i>REVISE</i>	-1187.61	2643.72	
vanilla		0.5	<i>ECCo</i>	-15.4966	17.1932	
vanilla		0.5	<i>Generic</i>	-11.7071	12.8003	
vanilla		0.5	<i>Omniscient</i>	-12.3897	14.1104	
vanilla		0.5	<i>REVISE</i>	-11.2965	13.1122	
full		1.0	<i>ECCo</i>	-6.06278	6.32519	
full		1.0	<i>Generic</i>	-1.57758e33	3.59342e33	
<b>full</b>	<b>1.0</b>	<b>Omniscient</b>	<b>-4.66436</b>	<b>4.88547</b>		
full		1.0	<i>REVISE</i>	-1157.26	2585.3	
vanilla		1.0	<i>ECCo</i>	-15.4915	17.2592	
vanilla		1.0	<i>Generic</i>	-10.8969	11.888	
vanilla		1.0	<i>Omniscient</i>	-12.6685	14.4499	
vanilla		1.0	<i>REVISE</i>	-11.2874	13.1369	
<b>full</b>	<b>5.0</b>	<b>ECCo</b>	<b>-2.56504</b>	<b>2.06543</b>		
full		5.0	<i>Generic</i>	-1.16971e28	2.66145e28	
full		5.0	<i>Omniscient</i>	-4.28955	4.30748	
full		5.0	<i>REVISE</i>	-530.204	1163.55	
vanilla		5.0	<i>ECCo</i>	-15.4763	17.1877	
vanilla		5.0	<i>Generic</i>	-11.6655	12.7364	
vanilla		5.0	<i>Omniscient</i>	-12.3937	14.1141	
vanilla		5.0	<i>REVISE</i>	-11.2976	13.0533	
<b>full</b>	<b>10.0</b>	<b>ECCo</b>	<b>-1.76439</b>	<b>0.973615</b>		
full		10.0	<i>Generic</i>	-1.54318e33	3.51163e33	
full		10.0	<i>Omniscient</i>	-4.44467	4.56008	
full		10.0	<i>REVISE</i>	-1515.03	3402.96	
vanilla		10.0	<i>ECCo</i>	-15.5074	17.275	
vanilla		10.0	<i>Generic</i>	-10.9077	11.8867	
vanilla		10.0	<i>Omniscient</i>	-12.6771	14.4225	
vanilla		10.0	<i>REVISE</i>	-11.2735	13.1031	
<b>full</b>	<b>15.0</b>	<b>ECCo</b>	<b>-1.36625</b>	<b>0.3652</b>		
full		15.0	<i>Generic</i>	-5.32108e28	1.21152e29	
full		15.0	<i>Omniscient</i>	-4.34376	4.38045	
full		15.0	<i>REVISE</i>	-473.027	1034.8	
vanilla		15.0	<i>ECCo</i>	-15.4703	17.1898	
vanilla		15.0	<i>Generic</i>	-11.6941	12.7669	
vanilla		15.0	<i>Omniscient</i>	-12.3895	14.0956	
vanilla		15.0	<i>REVISE</i>	-11.2868	13.0587	

**Table 3.** Results for Circles data by energy penalty.

Objective	Lambda	Energy (exper)	Generator	Type	Value	Std
full		<b>0.5</b>	ECCo	<b>-1.12388</b>	<b>0.216889</b>	
full		0.5	Generic	-1.20782	0.352005	
full		0.5	Omniscient	-5.09228	5.1182	
full		0.5	REVISE	-5.97244e27	1.36572e28	
vanilla		0.5	ECCo	-9.35338	7.30155	
vanilla		0.5	Generic	-8.89415	6.91671	
vanilla		0.5	Omniscient	-8.67963	6.9307	
vanilla		0.5	REVISE	-8.52507	6.74796	
full		<b>1.0</b>	ECCo	<b>-1.099</b>	<b>0.163365</b>	
full		1.0	Generic	-1.49485	0.726287	
full		1.0	Omniscient	-5.15975	5.20449	
full		1.0	REVISE	-3.09069e26	7.22344e26	
vanilla		1.0	ECCo	-9.33801	7.36386	
vanilla		1.0	Generic	-8.8619	6.85196	
vanilla		1.0	Omniscient	-8.69785	6.89941	
vanilla		1.0	REVISE	-8.69498	6.85371	
full		5.0	ECCo	-1.75204	0.154399	
full		<b>5.0</b>	Generic	<b>-1.21285</b>	<b>0.362686</b>	
full		5.0	Omniscient	-5.13516	5.16338	
full		5.0	REVISE	-1.09598e28	2.50339e28	
vanilla		5.0	ECCo	-9.36397	7.32382	
vanilla		5.0	Generic	-8.88498	6.90503	
vanilla		5.0	Omniscient	-8.70333	6.9289	
vanilla		5.0	REVISE	-8.51631	6.72565	
full		10.0	ECCo	-1.01708e6	2.31516e6	
full		<b>10.0</b>	Generic	<b>-1.48827</b>	<b>0.701741</b>	
full		10.0	Omniscient	-5.13432	5.15897	
full		10.0	REVISE	-3.74376e26	9.08858e26	
vanilla		10.0	ECCo	-9.31463	7.32684	
vanilla		10.0	Generic	-8.87348	6.86388	
vanilla		10.0	Omniscient	-8.7046	6.89274	
vanilla		10.0	REVISE	-8.68653	6.83497	
full		15.0	ECCo	-3.31332e13	7.53714e13	
full		<b>15.0</b>	Generic	<b>-1.21817</b>	<b>0.370377</b>	
full		15.0	Omniscient	-5.19548	5.23317	
full		15.0	REVISE	-9.01467e27	2.0592e28	
vanilla		15.0	ECCo	-9.37662	7.34277	
vanilla		15.0	Generic	-8.86149	6.8695	
vanilla		15.0	Omniscient	-8.69488	6.95691	
vanilla		15.0	REVISE	-8.50583	6.72685	